

# Objective - Image Classification

1. Solve image classification with **convolutional neural networks**.
2. Improve the performance with **data augmentations**.
3. Understand popular image model techniques such as **residual**.

# Task Introduction - Food Classification

- The images are collected from the food-11 dataset classified into 11 classes.
- Training set: 9866 labeled images
- Validation set: 3430 labeled images
- Testing set: 3347 images

# Rules

- DO NOT attempt to find the original labels of the testing set.
- DO NOT use any external datasets.
- **DO NOT use any pretrained models.**
  - Also, do not attempt to “test how effective pretraining is” by submitting to kaggle. Pretraining is very effective and you may test it after the competition ends.
- You may use any publicly available packages/code
  - But make sure you do not use pretrained models. Most code use those.
  - You may not upload your code/checkpoints to be publicly available during the timespan of this homework.

# Submission Format

The file should contain a header and have the following format:

```
Id,Category  
0001,1
```

Both type should be strings. Id corresponds to the jpg filenames in test. Follow the sample code if you have trouble with formatting.

# Model Selection

- Visit [torchvision.models](#) for a list of model structures, or go to [timm](#) for the latest model structures.
- Pretrained weights are not allowed, specifically set `pretrained=False` to ensure that the guideline is met.

## Classification

The models subpackage contains definitions for the following model architectures for image classification:

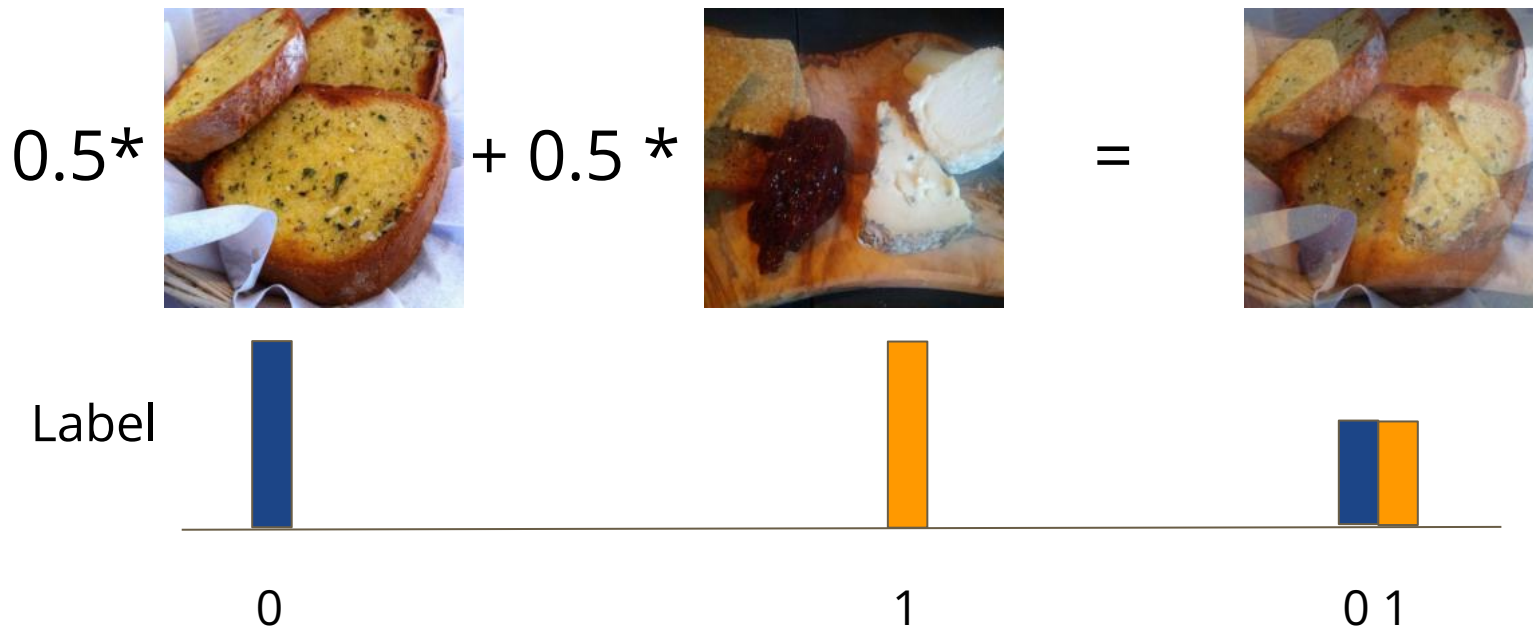
- AlexNet
- VGG
- ResNet
- SqueezeNet

# Data Augmentation

- Modify the image data so non-identical inputs are given to the model each epoch, to prevent overfitting of the model
- Visit [torchvision.transforms](https://pytorch.org/docs/stable/torchvision/transforms.html) for a list of choices and their corresponding effect. Diversity is encouraged! Usually, stacking multiple transformations leads to better results.
- Coding : fill in `train_tfm` to gain this effect



# Advanced Data Augmentation - mixup




# Advanced Data Augmentation - mixup

- Coding :
- In your `torch.utils.Dataset`, `__getitem__()` needs to return an image that is the linear combination of two images.
- In your `torch.utils.Dataset`, `__getitem__()` needs to return a label that is a vector, to assign probabilities to each class.
- You need to explicitly code out the math formula of the cross entropy loss, as `CrossEntropyLoss` does not support multiple labels.




# Test Time Augmentation

- The sample code tests images using a deterministic “test transformation”
- You may using the train transformation for a more diversified representation of the images, and predict with multiple variants of the test images.
- Coding : You need to fill in `train_tfm`, change the augmentation method for `test_dataset`, and modify prediction code to gain this effect



# Test Time Augmentation

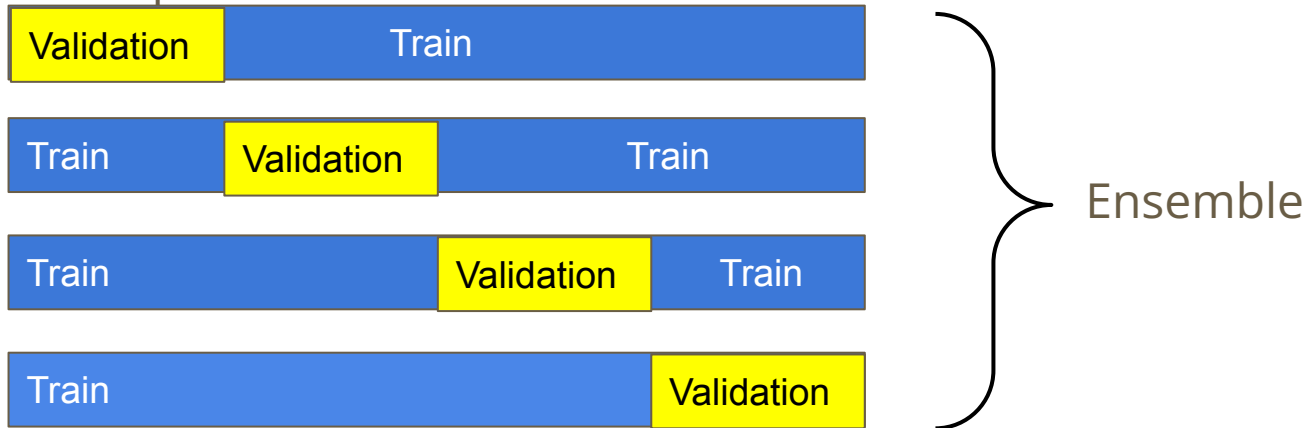
- Usually, test\_tfm will produce images that are more identifiable, so you can assign a larger weight to test\_tfm results for better performance.



- Ex : Final Prediction =  $\text{avg\_train\_tfm\_pred} * 0.5 + \text{test\_tfm\_pred} * 0.5$

# Cross Validation

- Cross-validation is a resampling method that uses different portions of the data to validate and train a model on different iterations. Ensembling multiple results lead to better performance.
- Coding : You need to merge the current train and validation paths, and resample from those to form new train and validation sets.



# Cross Validation

- Even if you don't do cross validation, you are encouraged to resplit the train/validation set to suitable proportions.
  - Currently, train : validation  $\sim 3 : 1$ , more training data could be valuable.